

CS390-NIP - Fall 2020
Name: Michael Lee
Email: lee2965@purdue.edu
Date: November 3, 2020

Lab 4 Report

Basic Image Classification

Due Date: November 2, 2020

<https://github.com/michaeldlee23/cs390-lab4>

Resources Used

- Conv2DTranspose and UpSampling2D in Keras
<https://machinelearningmastery.com/upsampling-and-transpose-convolution-layers-for-generative-adversarial-networks/>
- Calculating output size after Conv2DTranspose
<https://datascience.stackexchange.com/questions/26451/how-to-calculate-the-output-shape-of-conv2d-transpose>
- Tensorboard with train_on_batch
<https://stackoverflow.com/questions/44861149/keras-use-tensorboard-with-train-on-batch>

Completed Portions of Lab

1. [80] GANs
 - ✓[50] Complete the GAN to generate legible F-MNIST records. Generate from 3 classes.
 - ✓[10] Use a convnet for the GAN networks.
 - ✓[10] Implement the option to select the ratio of training between discriminator and generator.
 - ✓[10] Save a plot of loss over training steps for each network.
 - ✓? [+5] Generate very detailed F-MNIST records (there is no metric for this)
 - ✓? [+20] Generate legible CIFAR records from 3 classes
 - [+10] Create your own dataset of images and generate records of it. Images should be 32x32x1 or 32x32x3 and must be reasonably complex. Include the dataset in your submission with examples in your report.
2. ✓[20] Report

Questions

1. Describe the discriminator and generator.

The discriminator and generator are two parts of the GAN that basically train and improve each other — the generator produces images, and the discriminator takes those generated images, along with real images, and tries to classify a given image as real or generated. The generator tries to fool the discriminator, and the discriminator tries to learn how to not be tricked. Thus, we create an adversarial relationship between the two networks which results in an effective image generator.

2. Why do we sometimes need to train the discriminator and generator different amounts?

Due to the adversarial nature of the networks, we need to find the proper balance between the discriminator and generator in terms of their relative strength to each other. That is, if the discriminator is too good at its job, the generator won't gain enough information to improve. Likewise, if the generator is too good at its job, the discriminator will struggle to improve. One remedy to this is to adjust the training ratio between the two networks, allowing, for example, the generator to continue training while the discriminator stays constant until it can "catch up." We could achieve this by using a static constant that lets the generator get updates 10 times more frequently than the discriminator, for example, or we could set a threshold that only trains the discriminator when its loss reaches a certain level. Likewise, the same procedure could be applied to the generator.

Hyperparameters

Discriminator Model for MNIST data

Model: "MNIST_DISCRIMINATOR"

Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	(None, 14, 14, 64)	1664
leaky_re_lu_19 (LeakyReLU)	(None, 14, 14, 64)	0
dropout_15 (Dropout)	(None, 14, 14, 64)	0
conv2d_11 (Conv2D)	(None, 7, 7, 128)	73856
leaky_re_lu_20 (LeakyReLU)	(None, 7, 7, 128)	0
dropout_16 (Dropout)	(None, 7, 7, 128)	0
flatten_5 (Flatten)	(None, 6272)	0
dense_12 (Dense)	(None, 1024)	6423552
leaky_re_lu_21 (LeakyReLU)	(None, 1024)	0
dropout_17 (Dropout)	(None, 1024)	0
dense_13 (Dense)	(None, 1)	1025
Total params: 6,500,097		
Trainable params: 6,500,097		
Non-trainable params: 0		

Generator Model for MNIST data

Model: "MNIST_GENERATOR"

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 6272)	633472
leaky_re_lu_22 (LeakyReLU)	(None, 6272)	0
batch_normalization_4 (Batch Normalization)	(None, 6272)	25088
reshape_2 (Reshape)	(None, 7, 7, 128)	0
conv2d_transpose_4 (Conv2DTranspose)	(None, 14, 14, 64)	204864
leaky_re_lu_23 (LeakyReLU)	(None, 14, 14, 64)	0
batch_normalization_5 (Batch Normalization)	(None, 14, 14, 64)	256
conv2d_transpose_5 (Conv2DTranspose)	(None, 28, 28, 1)	1601
Total params: 865,281		
Trainable params: 852,609		
Non-trainable params: 12,672		

Discriminator Model for CIFAR data

Model: "CIFAR_DISCRIMINATOR"

Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	(None, 16, 16, 64)	1792
leaky_re_lu_16 (LeakyReLU)	(None, 16, 16, 64)	0
conv2d_11 (Conv2D)	(None, 8, 8, 128)	73856
leaky_re_lu_17 (LeakyReLU)	(None, 8, 8, 128)	0
conv2d_12 (Conv2D)	(None, 4, 4, 128)	147584
leaky_re_lu_18 (LeakyReLU)	(None, 4, 4, 128)	0
conv2d_13 (Conv2D)	(None, 2, 2, 256)	295168
leaky_re_lu_19 (LeakyReLU)	(None, 2, 2, 256)	0
flatten_2 (Flatten)	(None, 1024)	0
dropout_2 (Dropout)	(None, 1024)	0
dense_4 (Dense)	(None, 1)	1025
Total params: 519,425		
Trainable params: 519,425		
Non-trainable params: 0		

Generator Model for CIFAR data

Model: "CIFAR_GENERATOR"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 4096)	413696
leaky_re_lu_20 (LeakyReLU)	(None, 4096)	0
reshape_2 (Reshape)	(None, 4, 4, 256)	0
conv2d_transpose_6 (Conv2DTr	(None, 8, 8, 128)	524416
leaky_re_lu_21 (LeakyReLU)	(None, 8, 8, 128)	0
conv2d_transpose_7 (Conv2DTr	(None, 16, 16, 128)	262272
leaky_re_lu_22 (LeakyReLU)	(None, 16, 16, 128)	0
conv2d_transpose_8 (Conv2DTr	(None, 32, 32, 128)	262272
leaky_re_lu_23 (LeakyReLU)	(None, 32, 32, 128)	0
conv2d_14 (Conv2D)	(None, 32, 32, 3)	6147
Total params: 1,468,803		
Trainable params: 1,468,803		
Non-trainable params: 0		

Relevant Plots - MNIST_F




Figure 1: Loss over time for MNIST_F Sneaker GAN




Figure 2: Loss over time for MNIST_F Coat GAN




Figure 3: Loss over time for MNIST_F Trouser GAN

Generated Images - MNIST_F




Figure 4: Generated Sneaker Images



Figure 5: Generated Coat Images




Figure 6: Generated Trouser Images

Relevant Plots - CIFAR_10




Figure 7: Loss over time for CIFAR_10 Horse GAN

Interestingly, while there is a clear divergence in loss which would suggest that the GAN is probably not performing well, we do nevertheless get a couple horse-like images. We have similar results for Bird and Dog.

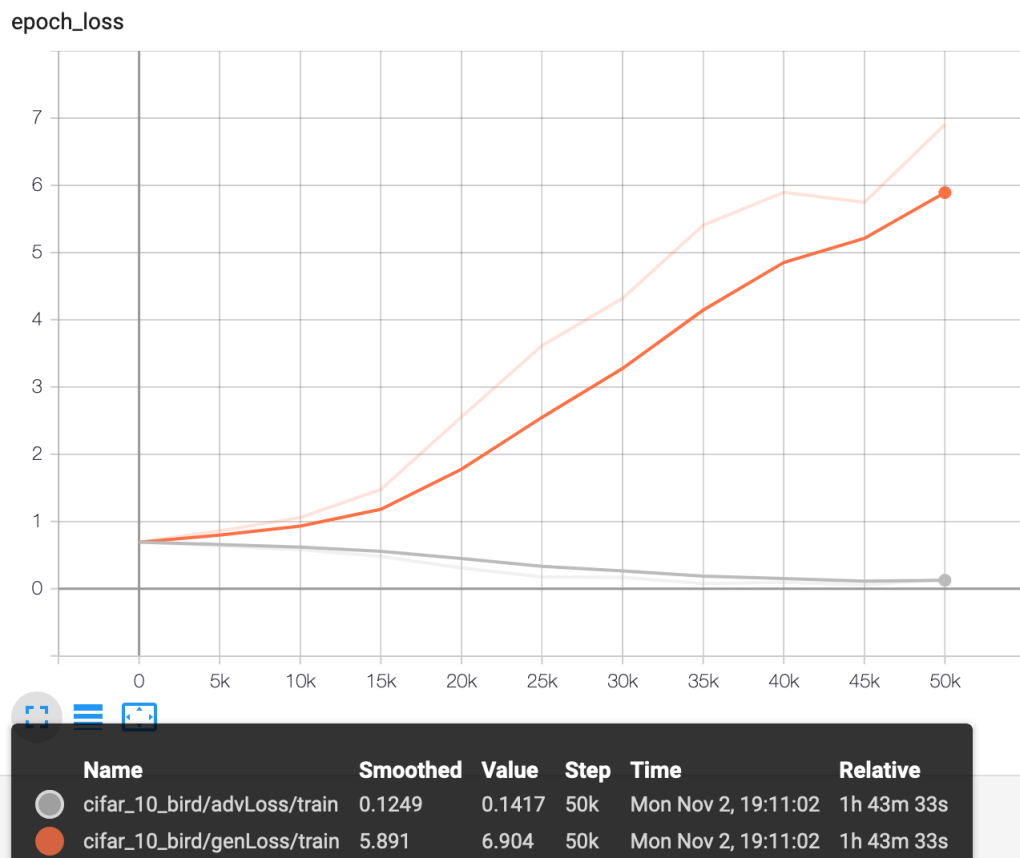


Figure 8: Loss over time for CIFAR_10 Bird GAN

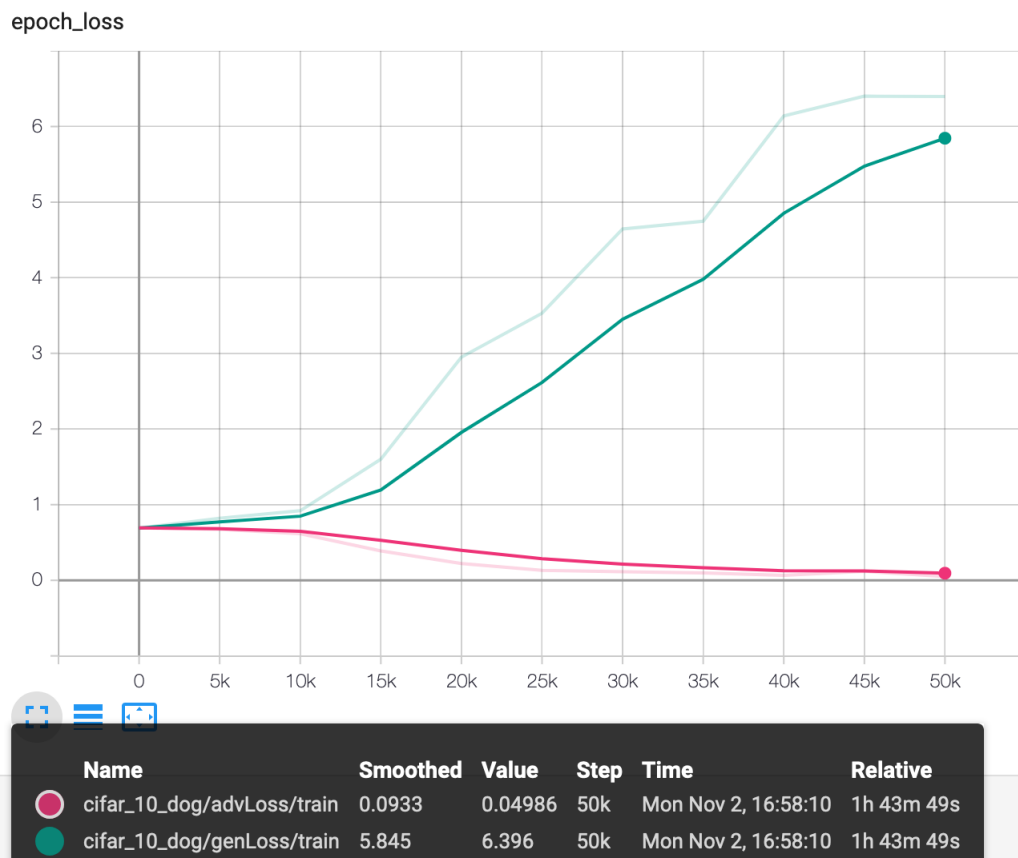


Figure 9: Loss over time for CIFAT_10 Dog GAN

Generated Images - CIFAR_10




Figure 10: Generated Horse Images

While some of these pictures are pretty illegible, we still get a couple that definitely have a horse-like quality to them. Namely, the last two images in the first row.




Figure 11: Generated Bird Images

The bird generator did not do that well, although the very last image does sort of look like a bird. Note that the bird generator may have produced better pictures, similar to the horse generator, if I trained it longer. The horse generator was trained twice as long as the bird and dog generator.




Figure 12: Generated Dog Images

The dog generator also did not perform very well, although we can kind of see the likeness of a dog in the very last picture. The right side of the image gets a little disfigured and what looks to be a tongue seems to be unnatural, but we we can definitely see an attempt at a dog.

Generated Images - CIFAR_10, Cherrypicked Horses




Figure 13: Cherrypicked Generated Horse Images

These were the best horse images after visually inspecting both the final and test outputs.